

## 6.824 2020 Lecture 4: Primary/Backup Replication

### Today

Primary/Backup Replication for Fault Tolerance  
Case study of VMware FT, an extreme version of the idea

The topic is (still) fault tolerance  
to provide availability  
despite server and network failures  
using replication

What kinds of failures can replication deal with?

"fail-stop" failure of a single replica  
fan stops working, CPU overheats and shuts itself down  
someone trips over replica's power cord or network cable  
software notices it is out of disk space and stops  
Maybe not defects in h/w or bugs in s/w or human configuration errors  
Often not fail-stop  
May be correlated (i.e. cause all replicas to crash at the same time)  
But, sometimes can be detected (e.g. checksums)  
How about earthquake or city-wide power failure?  
Only if replicas are physically separated

Is replication worth the Nx expense?

Two main replication approaches:

State transfer  
Primary replica executes the service  
Primary sends [new] state to backups  
Replicated state machine  
Clients send operations to primary,  
primary sequences and sends to backups  
All replicas execute all operations  
If same start state,  
same operations,  
same order,  
deterministic,  
then same end state.

State transfer is simpler

But state may be large, slow to transfer over network

Replicated state machine often generates less network traffic

Operations are often small compared to state  
But complex to get right  
VM-FT uses replicated state machine, as do Labs 2/3/4

Big Questions:

What state to replicate?  
Does primary have to wait for backup?  
When to cut over to backup?  
Are anomalies visible at cut-over?  
How to bring a replacement backup up to speed?

At what level do we want replicas to be identical?

Application state, e.g. a database's tables?  
GFS works this way  
Can be efficient; primary only sends high-level operations to backup  
Application code (server) must understand fault tolerance, to e.g. forward op stream  
Machine level, e.g. registers and RAM content?  
might allow us to replicate any existing server w/o modification!  
requires forwarding of machine events (interrupts, DMA, &c)  
requires "machine" modifications to send/recv event stream...

Today's paper (VMware FT) replicates machine-level state

Transparent: can run any existing O/S and server software!  
Appears like a single server to clients

Overview

[diagram: app, O/S, VM-FT underneath, disk server, network, clients]  
words:

hypervisor == monitor == VMM (virtual machine monitor)  
O/S+app is the "guest" running inside a virtual machine  
two machines, primary and backup  
primary sends all external events (client packets &c) to backup over network  
"logging channel", carrying log entries  
ordinarily, backup's output is suppressed by FT  
if either stops being able to talk to the other over the network  
"goes live" and provides sole service  
if primary goes live, it stops sending log entries to the backup

VMM emulates a local disk interface  
but actual storage is on a network server  
treated much like a client:  
usually only primary communicates with disk server (backup's FT discards)  
if backup goes live, it talks to disk server  
external disk makes creating a new backup faster (don't have to copy primary's disk)

When does the primary have to send information to the backup?  
Any time something happens that might cause their executions to diverge.  
Anything that's not a deterministic consequence of executing instructions.

What sources of divergence must FT handle?  
Most instructions execute identically on primary and backup.  
As long as memory+registers are identical,  
which we're assuming by induction.  
Inputs from external world -- just network packets.  
These appear as DMA'd data plus an interrupt.  
Timing of interrupts.  
Instructions that aren't functions of state, such as reading current time.  
Not multi-core races, since uniprocessor only.

Why would divergence be a disaster?  
b/c state on backup would differ from state on primary,  
and if primary then failed, clients would see inconsistency.  
Example: GFS lease expiration  
Imagine we're replicating the GFS master  
Chunkserver must send "please renew" msg before 60-second lease expires  
Clock interrupt drives master's notion of time  
Suppose chunkserver sends "please renew" just around 60 seconds  
On primary, clock interrupt happens just after request arrives.  
Primary copy of master renews the lease, to the same chunkserver.  
On backup, clock interrupt happens just before request.  
Backup copy of master expires the lease.  
If primary fails, backup takes over, it will think there  
is no lease, and grant it to a different chunkserver.  
Then two chunkservers will have lease for same chunk.  
So: backup must see same events,  
in same order,  
at same points in instruction stream.

Each log entry: instruction #, type, data.

FT's handling of timer interrupts  
Goal: primary and backup should see interrupt at  
the same point in the instruction stream  
Primary:  
FT fields the timer interrupt  
FT reads instruction number from CPU  
FT sends "timer interrupt at instruction X" on logging channel  
FT delivers interrupt to primary, and resumes it  
(this relies on CPU support to interrupt after the X'th instruction)  
Backup:  
ignores its own timer hardware  
FT sees log entry \*before\* backup gets to instruction X  
FT tells CPU to interrupt (to FT) at instruction X  
FT mimics a timer interrupt to backup

## FT's handling of network packet arrival (input)

### Primary:

FT tells NIC to copy packet data into FT's private "bounce buffer"  
 At some point NIC does DMA, then interrupts  
 FT gets the interrupt  
 FT pauses the primary  
 FT copies the bounce buffer into the primary's memory  
 FT simulates a NIC interrupt in primary  
 FT sends the packet data and the instruction # to the backup

### Backup:

FT gets data and instruction # from log stream  
 FT tells CPU to interrupt (to FT) at instruction X  
 FT copies the data to backup memory, simulates NIC interrupt in backup

## Why the bounce buffer?

We want the data to appear in memory at exactly the same point in execution of the primary and backup.  
 Otherwise they may diverge.

## Note that the backup must lag by one one log entry

Suppose primary gets an interrupt, or input, after instruction X  
 If backup has already executed past X, it cannot handle the input correctly  
 So backup FT can't start executing at all until it sees the first log entry  
 Then it executes just to the instruction # in that log entry  
 And waits for the next log entry before resuming backup

## Example: non-deterministic instructions

some instructions yield different results even if primary/backup have same state  
 e.g. reading the current time or cycle count or processor serial #

### Primary:

FT sets up the CPU to interrupt if primary executes such an instruction  
 FT executes the instruction and records the result  
 sends result and instruction # to backup

### Backup:

FT reads log entry, sets up for interrupt at instruction #  
 FT then supplies value that the primary got

## What about output (sending network packets)?

Primary and backup both execute instructions for output  
 Primary's FT actually does the output  
 Backup's FT discards the output

## Output example: DB server

clients can send "increment" request  
 DB increments stored value, replies with new value  
 so:

[diagram]

suppose the server's value starts out at 10  
 network delivers client request to FT on primary  
 primary's FT sends on logging channel to backup  
 FTs deliver request to primary and backup  
 primary executes, sets value to 11, sends "11" reply, FT really sends reply  
 backup executes, sets value to 11, sends "11" reply, and FT discards  
 the client gets one "11" response, as expected

## But wait:

suppose primary crashes just after sending the reply  
 so client got the "11" reply  
 AND the logging channel discards the log entry w/ client request  
 primary is dead, so it won't re-send  
 backup goes live  
 but it has value "10" in its memory!  
 now a client sends another increment request  
 it will get "11" again, not "12"  
 oops

## Solution: the Output Rule (Section 2.2)

before primary sends output,  
 must wait for backup to acknowledge all previous log entries

Again, with output rule:

[diagram]

primary:

- receives client "increment" request
- sends client request on logging channel
- about to send "11" reply to client
- first waits for backup to acknowledge previous log entry
- then sends "11" reply to client

suppose the primary crashes at some point in this sequence

if before primary receives acknowledgement from backup

- maybe backup didn't see client's request, and didn't increment
- but also primary won't have replied

if after primary receives acknowledgement from backup

- then client may see "11" reply
- but backup guaranteed to have received log entry w/ client's request
- so backup will increment to 11

The Output Rule is a big deal

- Occurs in some form in all replication systems

- A serious constraint on performance

- An area for application-specific cleverness

- Eg. maybe no need for primary to wait before replying to read-only operation

- FT has no application-level knowledge, must be conservative

Q: What if the primary crashes just after getting ACK from backup,  
but before the primary emits the output?

Does this mean that the output won't ever be generated?

A: Here's what happens when the primary fails and the backup goes live.

- The backup got some log entries from the primary.

- The backup continues executing those log entries WITH OUTPUT DISCARDED.

- After the last log entry, the backup goes live -- stops discarding output

- In our example, the last log entry is arrival of client request

- So after client request arrives, the client will start emitting outputs

- And thus it will emit the reply to the client

Q: But what if the primary crashed *after* emitting the output?

Will the backup emit the output a *second* time?

A: Yes.

- OK for TCP, since receivers ignore duplicate sequence numbers.

- OK for writes to disk, since backup will write same data to same block #.

Duplicate output at cut-over is pretty common in replication systems

- Clients need to keep enough state to ignore duplicates

- Or be designed so that duplicates are harmless

Q: Does FT cope with network partition -- could it suffer from split brain?

- E.g. if primary and backup both think the other is down.

- Will they both go live?

A: The disk server breaks the tie.

- Disk server supports atomic test-and-set.

- If primary or backup thinks other is dead, attempts test-and-set.

- If only one is alive, it will win test-and-set and go live.

- If both try, one will lose, and halt.

The disk server may be a single point of failure

- If disk server is down, service is down

- They probably have in mind a replicated disk server

Q: Why don't they support multi-core?

Performance (table 1)

- FT/Non-FT: impressive!

- little slow down

- Logging bandwidth

- Directly reflects disk read rate + network input rate

18 Mbit/s for my-sql  
These numbers seem low to me  
Applications can read a disk at at least 400 megabits/second  
So their applications aren't very disk-intensive

When might FT be attractive?

Critical but low-intensity services, e.g. name server.  
Services whose software is not convenient to modify.

What about replication for high-throughput services?

People use application-level replicated state machines for e.g. databases.  
The state is just the DB, not all of memory+disk.  
The events are DB commands (put or get), not packets and interrupts.  
Result: less fine-grained synchronization, less overhead.  
GFS use application-level replication, as do Lab 2 &c

Summary:

Primary-backup replication  
VM-FT: clean example  
How to cope with partition without single point of failure?  
Next lecture  
How to get better performance?  
Application-level replicated state machines

----

VMware KB (#1013428) talks about multi-CPU support. VM-FT may have switched from a replicated state machine approach to the state transfer approach, but unclear whether that is true or not.

<http://www.wooditwork.com/2014/08/26/whats-new-vsphere-6-0-fault-tolerance/>

<http://www-mount.ece.umn.edu/~jjyi/MoBS/2007/program/01C-Xu.pdf>